Master Tehnologii Informatice
Administrarea Bazelor de Date
Anul universitar 2023-2024

Universitatea Politehnica din Timișoara

Administarea Bazelor de Date

-Proiect MongoDB-

Masterand: Mateișescu Maria
Prof. Coordonator: Bocan Valer

Sesiunea: Iunie 2024

Overview

The goal of the project is to assess the MongoDB database skills of the student. The project has a medium difficulty level and is relevant to industry employers of today.

Project description

You are required to make several statistical computations on some US Zips dataset using MongoDB as the database platform.

Prerequisites:

- Download the latest US Zips dataset from (choose the free tier). The dataset has approximately 33k entries.
- Create a MongoDB instance. You may use your own MongoDB Atlas instance in cloud or use a local instance. For local instances Docker is preferred, but you may also choose to install MongoDB as a standalone server on your OS.
- Import the dataset into the MongoDB instance.

Requirements

a) Get the states with a total population of over 10 million.

b) Get the average city population by state.

c) Get the largest and the smallest city in each state.

d) Get the largest and the smallest counties in each state.

e) Get the nearest 10 zips from one of Chicago's landmarks, the Willis Tower situated at coordinates 41.878876, -87.635918.

f) Get the total population situated between 50 and 200 kms around New York's landmark, the Statue of Liberty at coordinates 40.689247, -74.044502.

Primul pas a fost instalarea utilitarului MongoShell si a uneltelor MongoDB Database Tools pentru sistemul meu de operare (Windows 10). Ulterior, am creat un cluster in MongoDB Atlas, in proiectul denumit „ABD project”.

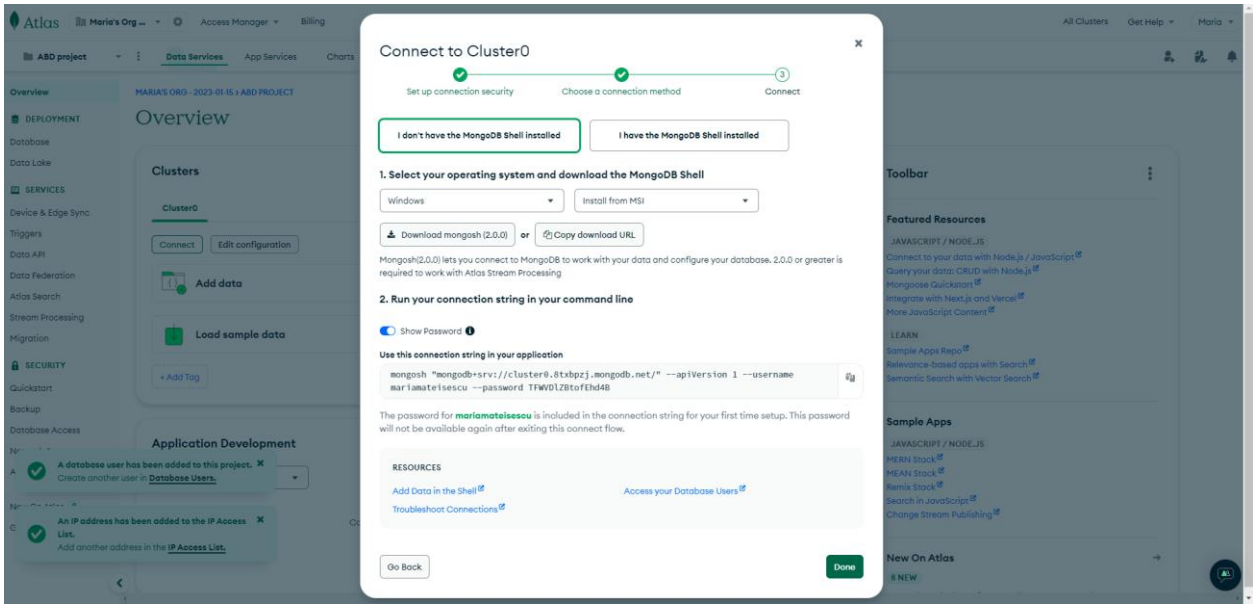

Pentru a ma conecta la instanta din cloud MongoDB Atlas am folosit connection string-ul obtinut anterior, care contine username-ul si parola.

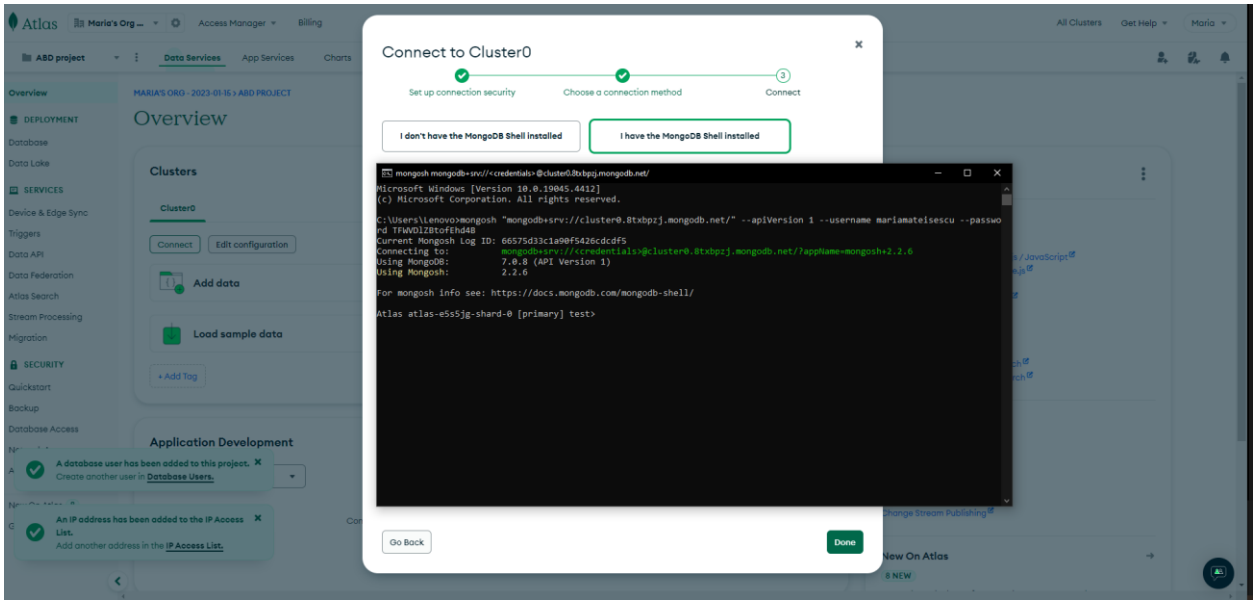

Unitatea de stocare a bazelor de date non-relationale este colectia, deci in baza de date vor fi date stocate sub forma de colectii, spre deosebire de bazele de date relationale unde exista tabele.

Folosind comanda „mongoimport” pusa la dispozitie de tool-uri am importat fisierul denumit „uszips” cu extensia .csv in baza de date „test”, iar colectia am denumit-o „uszips”. In total s-au importat cu succes 33787 de documente.

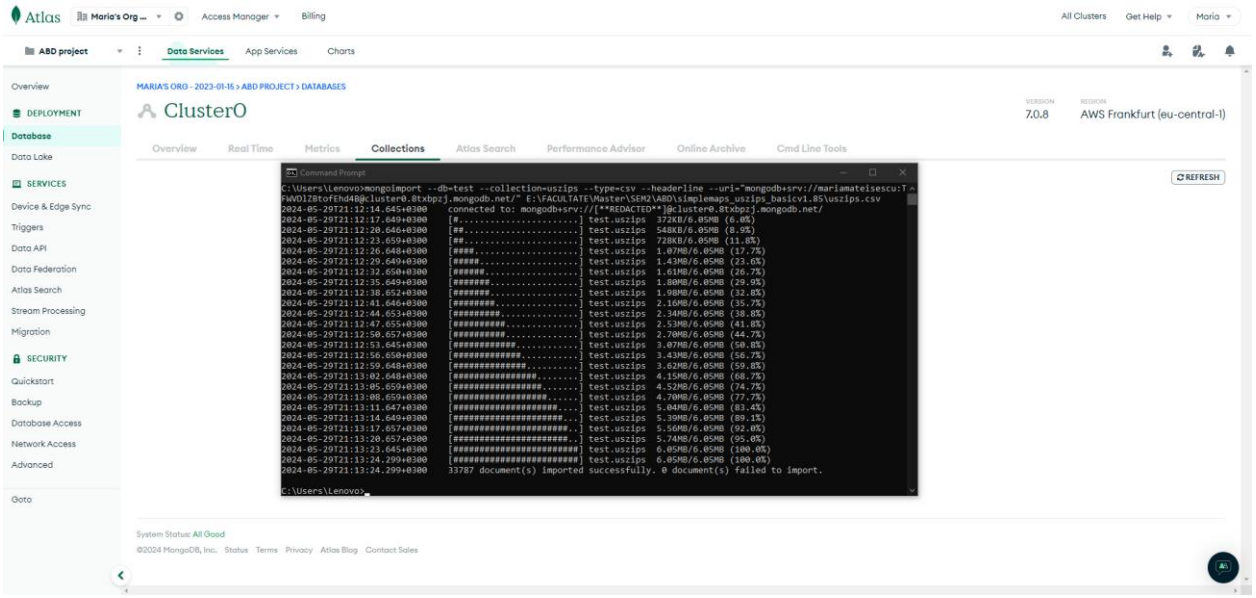

Dupa finalizarea importului, documentele din colectie pot fi vizualizate structurat in MongoDB Atlas. Pentru fiecare document se genereaza in mod automat campul _id de tipul ObjectId, reprezentand identificatorul unic.

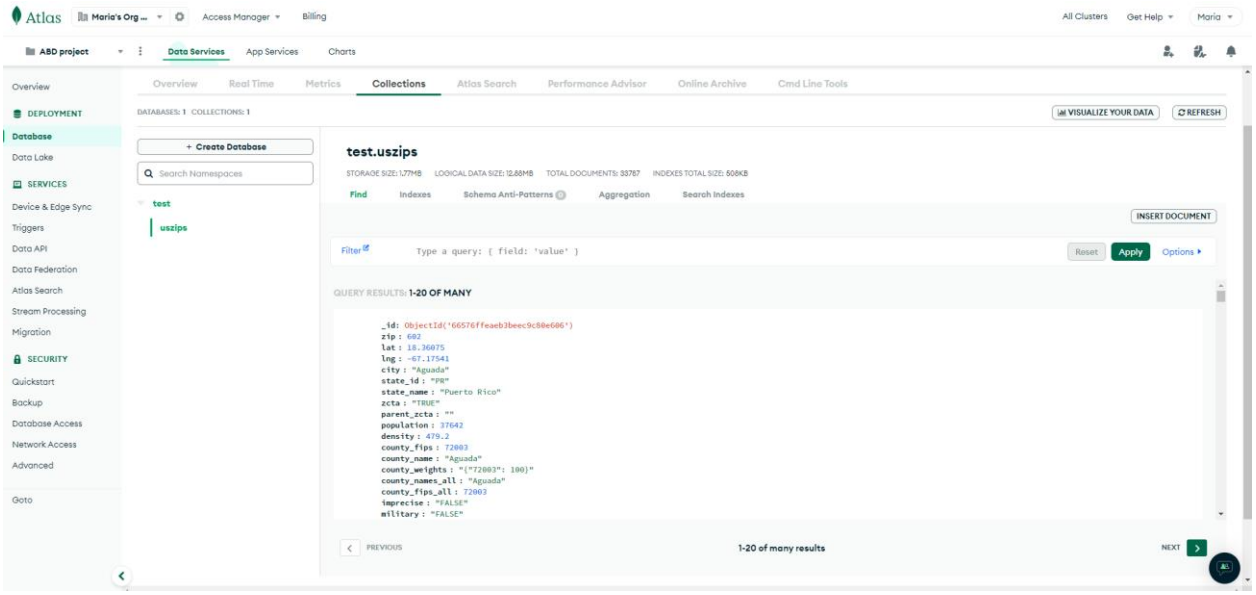

a) Obtineti statele cu o populatie totala de peste 10 milioane

Folosind comanda db.usizps.find().pretty() se obtin documentele intr-un format mai usor de citit. Pentru a obtine statele cu populatie totala de peste 10 milioane am folosit comanda:

db.uszips.aggregate([{ $group: {_id: "$state_name", total_population: { $sum: "$population" }}}, { $match: { total_population: { $gt: 10000000 }}}])

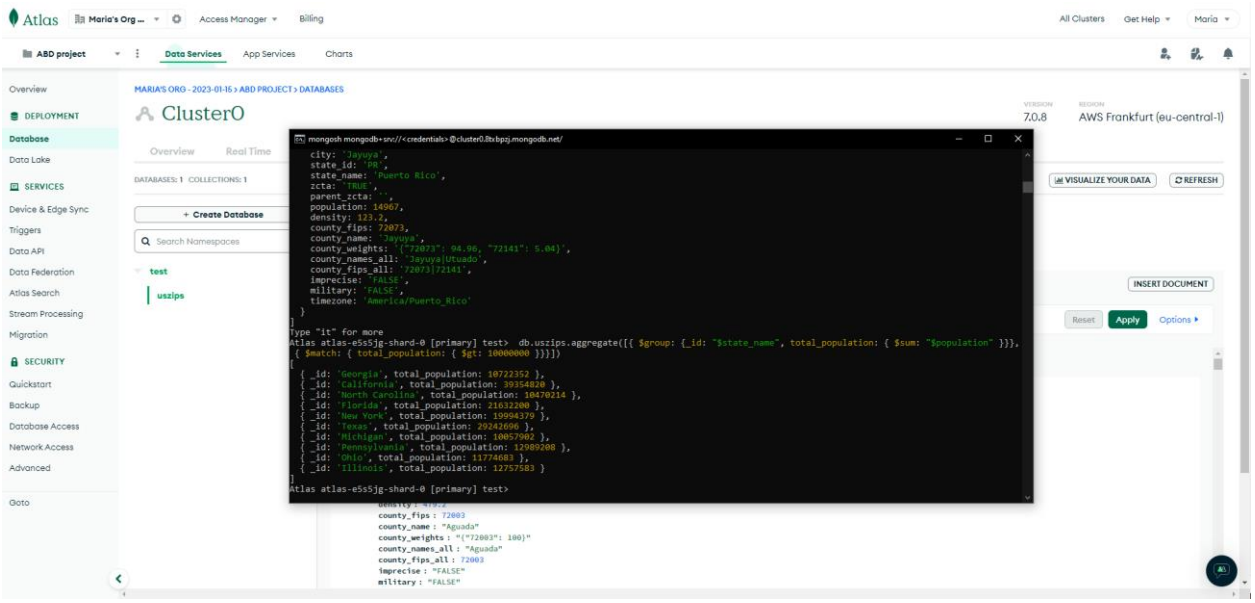

b) Obtineti populatia medie pe oras in functie de stat

Am folosit comanda:

db.uszips.aggregate([{ $group: { _id: { state: "$state_name", city: "$city" }, city_population: { $sum: "$population" } } }, { $group: { _id: "$_id.state", average_city_population: { $avg: "$city_population" } } }, { $project: { _id: 0, state_name: "$_id", average_city_population: 1 } }]).pretty()

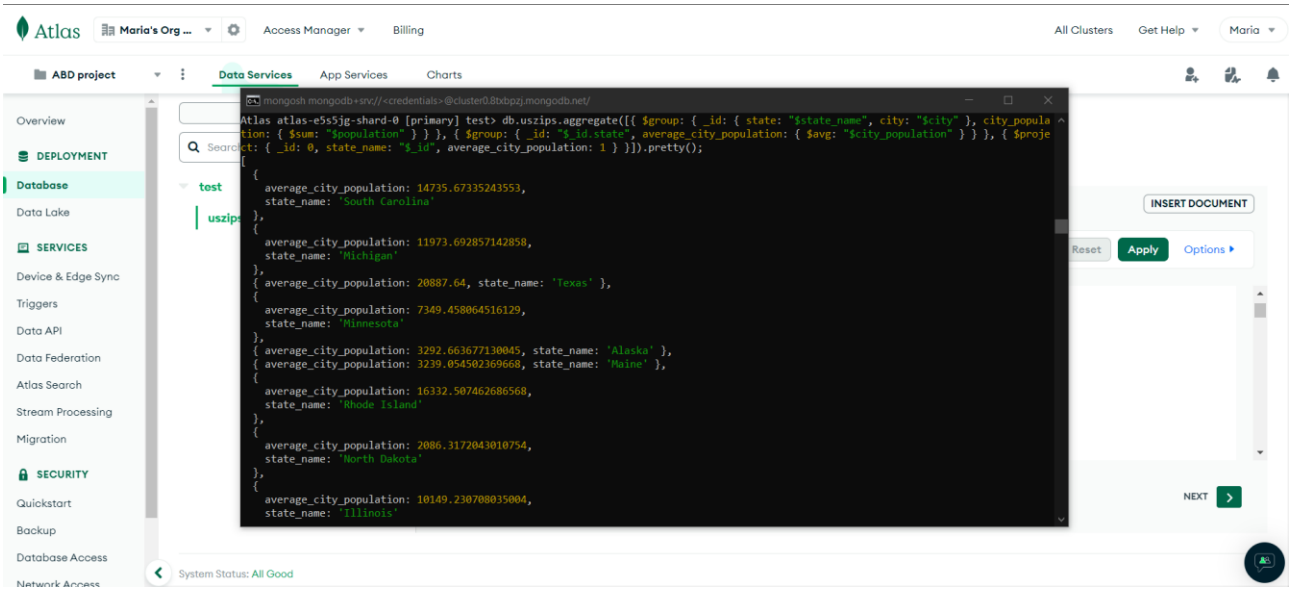

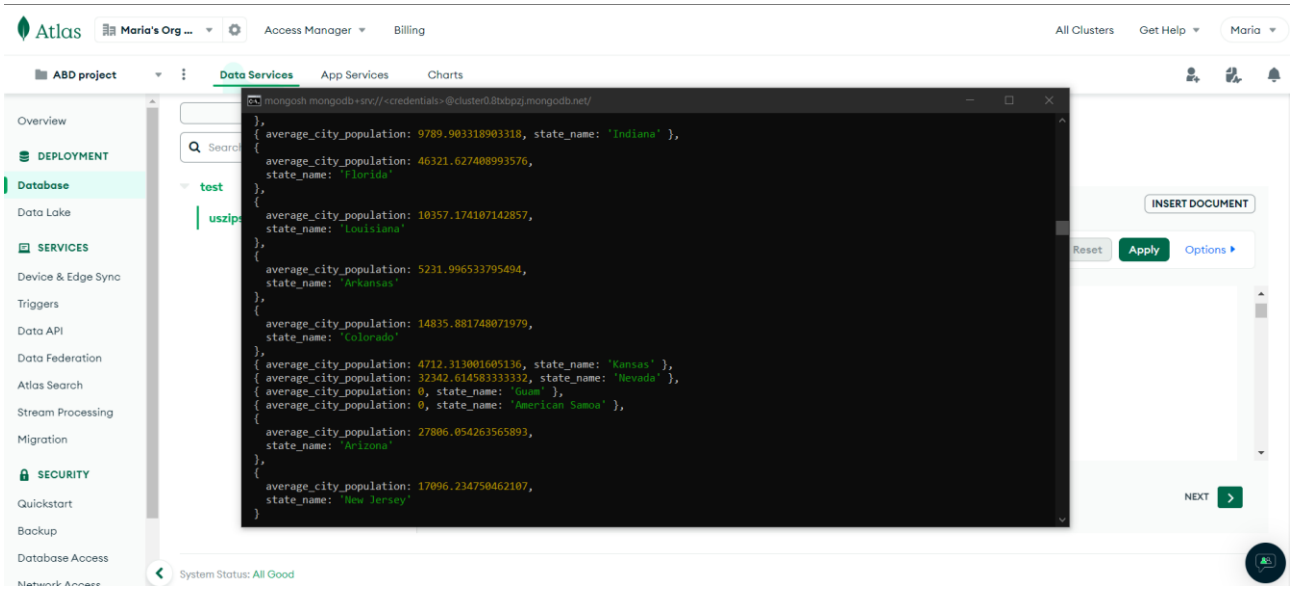
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
App Services
Charts
mongosh mongodb+srv://<credentials>@cluster0.8txbpzj.mongodb.net/
{ average_city_population: 9789.903318903318, state_name: 'Indiana' },
{ average_city_population: 46321.627408993576, state_name: 'Florida' },
{ average_city_population: 10357.174107142857, state_name: 'Louisiana' },
{ average_city_population: 5231.996533795494, state_name: 'Arkansas' },
{ average_city_population: 14835.881748071979, state_name: 'Colorado' },
{ average_city_population: 4712.313001605136, state_name: 'Kansas' },
{ average_city_population: 32342.614583333332, state_name: 'Nevada' },
{ average_city_population: 0, state_name: 'Guam' },
{ average_city_population: 0, state_name: 'American Samoa' },
{ average_city_population: 27806.054263565893, state_name: 'Arizona' },
{ average_city_population: 17096.234750462107, state_name: 'New Jersey' }
Overview
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
test
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good

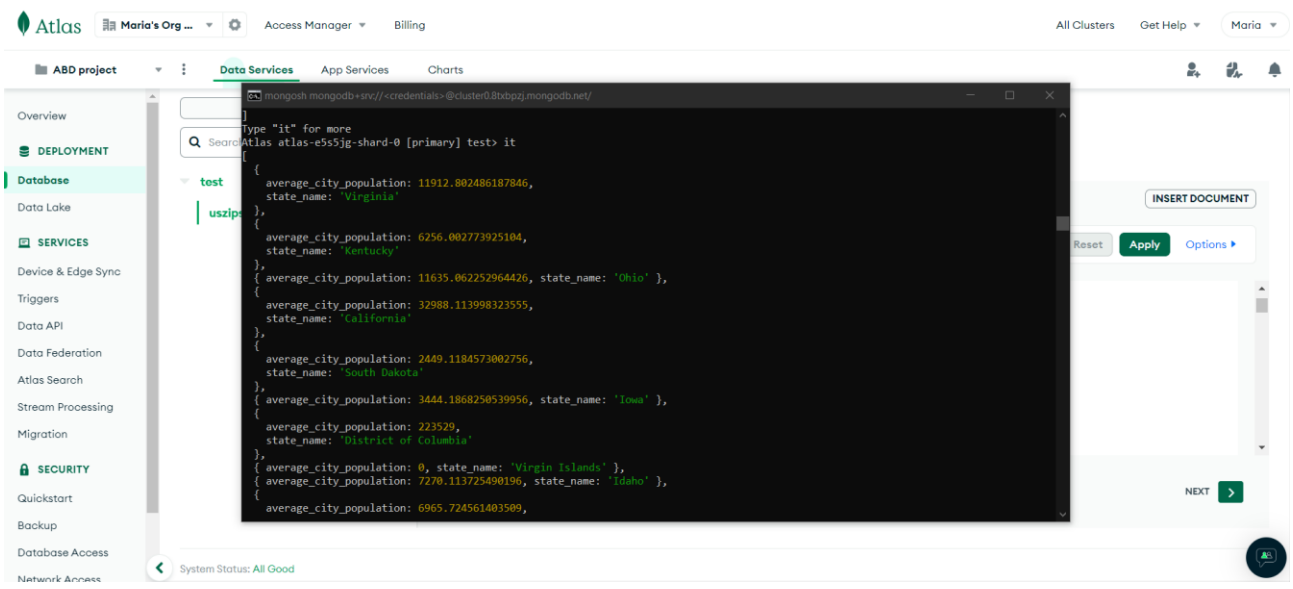
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
App Services
Charts
mongosh mongodb+srv://<credentials>@cluster0.8txbpzj.mongodb.net/
Type "it" for more
Atlas atlas-e5s5jg-shard-0 [primary] test> it
{ average_city_population: 11912.802486187846, state_name: 'Virginia' },
{ average_city_population: 6256.002773925104, state_name: 'Kentucky' },
{ average_city_population: 11635.062252964426, state_name: 'Ohio' },
{ average_city_population: 32988.113998323555, state_name: 'California' },
{ average_city_population: 2449.1184573002756, state_name: 'South Dakota' },
{ average_city_population: 3444.1868250539956, state_name: 'Iowa' },
{ average_city_population: 223529, state_name: 'District of Columbia' },
{ average_city_population: 0, state_name: 'Virgin Islands' },
{ average_city_population: 7270.113725490196, state_name: 'Idaho' },
{ average_city_population: 6965.724561403509,
Overview
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
test
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good

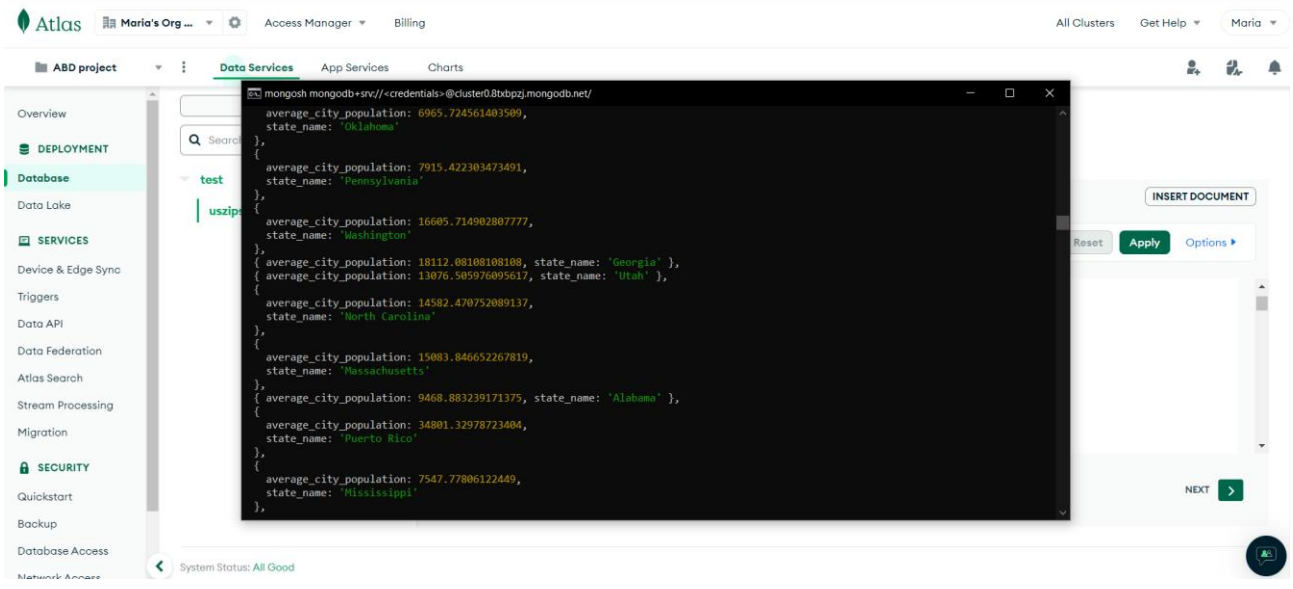
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
App Services
Charts
mongosh mongodb+srv://<credentials>@cluster0.8txbpzj.mongodb.net/
average_city_population: 6965.724561403509, state_name: 'Oklahoma' },
{ average_city_population: 7915.422303473491, state_name: 'Pennsylvania' },
{ average_city_population: 16605.714902807777, state_name: 'Washington' },
{ average_city_population: 18112.08108108108, state_name: 'Georgia' },
{ average_city_population: 13076.505976095617, state_name: 'Utah' },
{ average_city_population: 14582.470752089137, state_name: 'North Carolina' },
{ average_city_population: 15083.846652267819, state_name: 'Massachusetts' },
{ average_city_population: 9468.883239171375, state_name: 'Alabama' },
{ average_city_population: 34801.32978723404, state_name: 'Puerto Rico' },
{ average_city_population: 7547.77806122449, state_name: 'Mississippi' },
Overview
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
test
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good

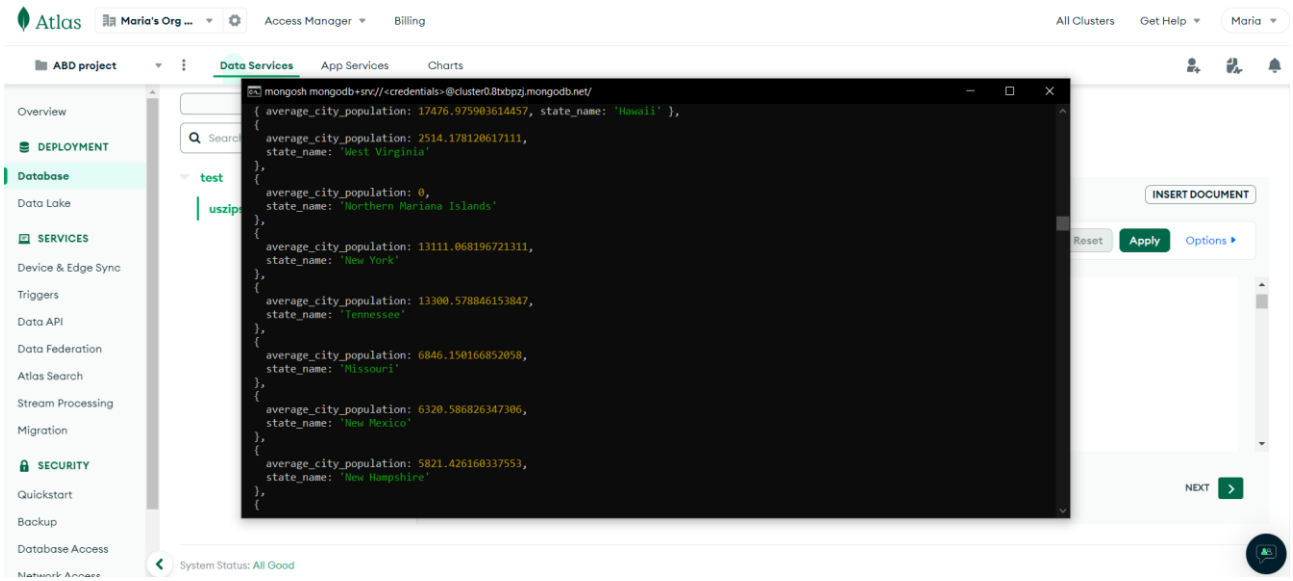

c) Obtineti cel mai mare si cel mai mic oras din fiecare stat

Am folosit comanda:

db.uszips.aggregate([{ $group: { _id: { state: "$state_name", city: "$city" }, city_population: { $sum: "$population" }}}, { $sort: { "_id.state": 1, city_population: -1 }}, { $group: { _id: "$_id.state", largest_city: { $first: { city: "$_id.city", population: "$city_population" }}, smallest_city: { $last: { city: "$_id.city", population: "$city_population" }}}}, { $project: { _id: 0, state_name: "$_id", largest_city: 1, smallest_city: 1 }}])

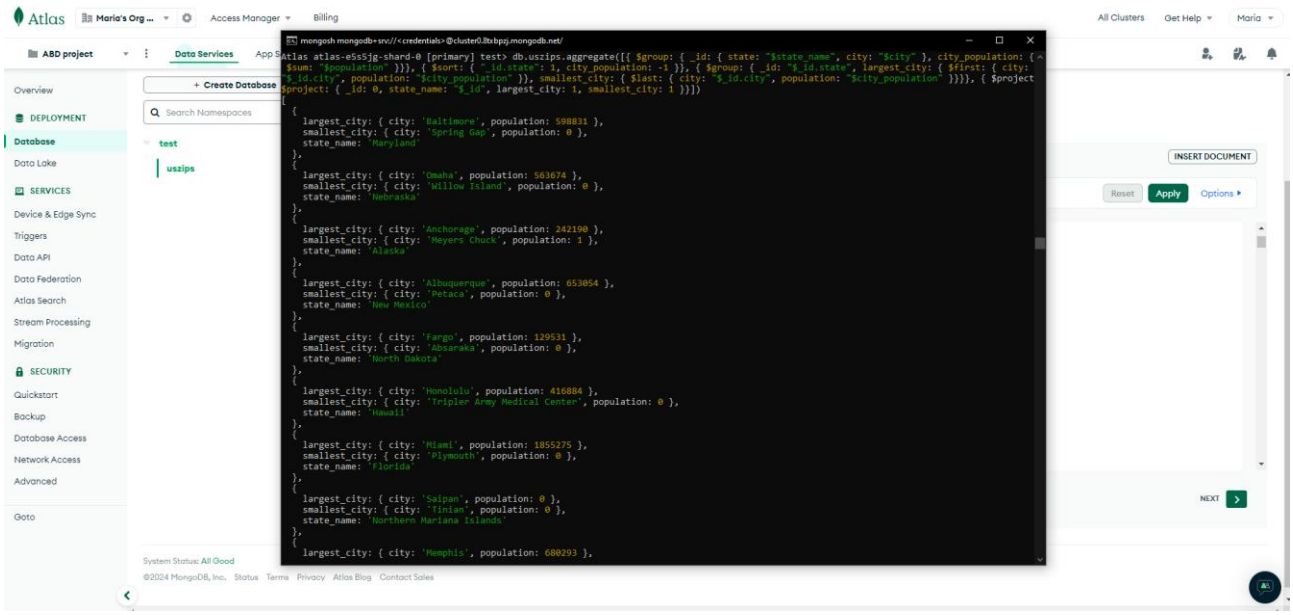

d) Obtineti cel mai mare si cel mai mic judet din fiecare stat

Am folosit comanda:

db.uszips.aggregate([{ $match: { population: { $exists: true, $ne: null } } }, { $group: { _id: { state: "$state_name", county: "$county_name" }, county_population: { $sum: "$population" } } }, { $sort: { "_id.state": 1, county_population: -1 } }, { $group: { _id: "$_id.state", largest_county: { $first: { county: "$_id.county", population: "$county_population" } }, smallest_county: { $last: { county: "$_id.county", population: "$county_population" } } } }, { $project: { _id: 0, state_name: "$_id", largest_county: 1, smallest_county: 1 } }])

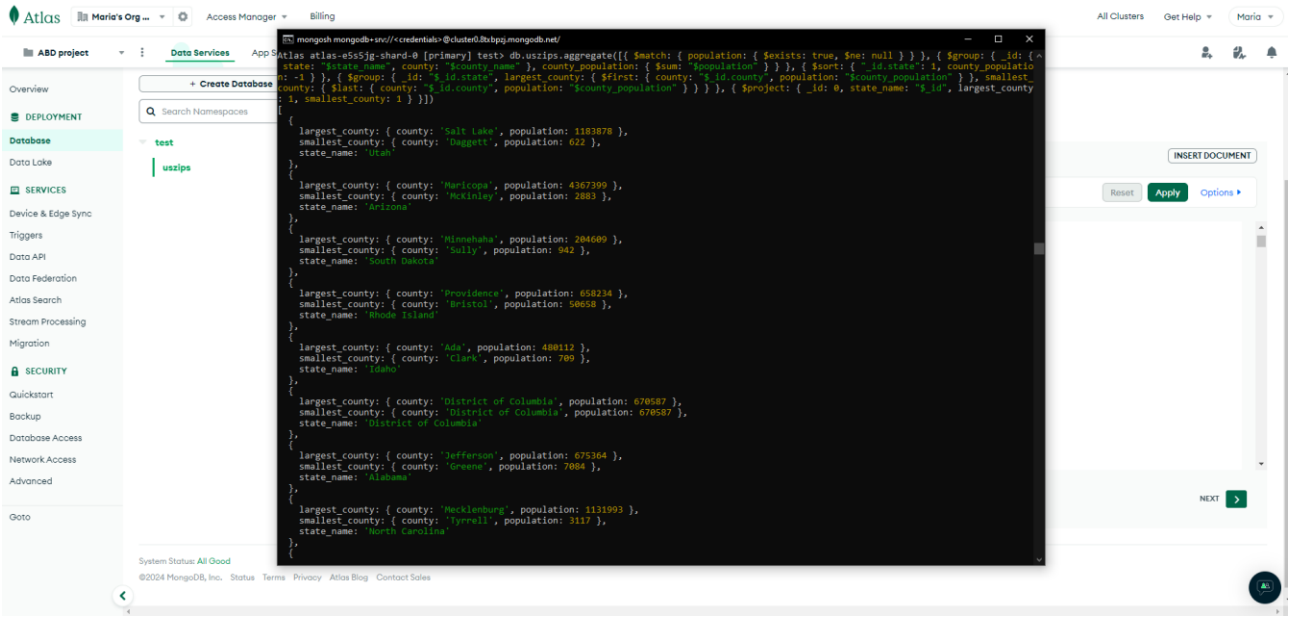

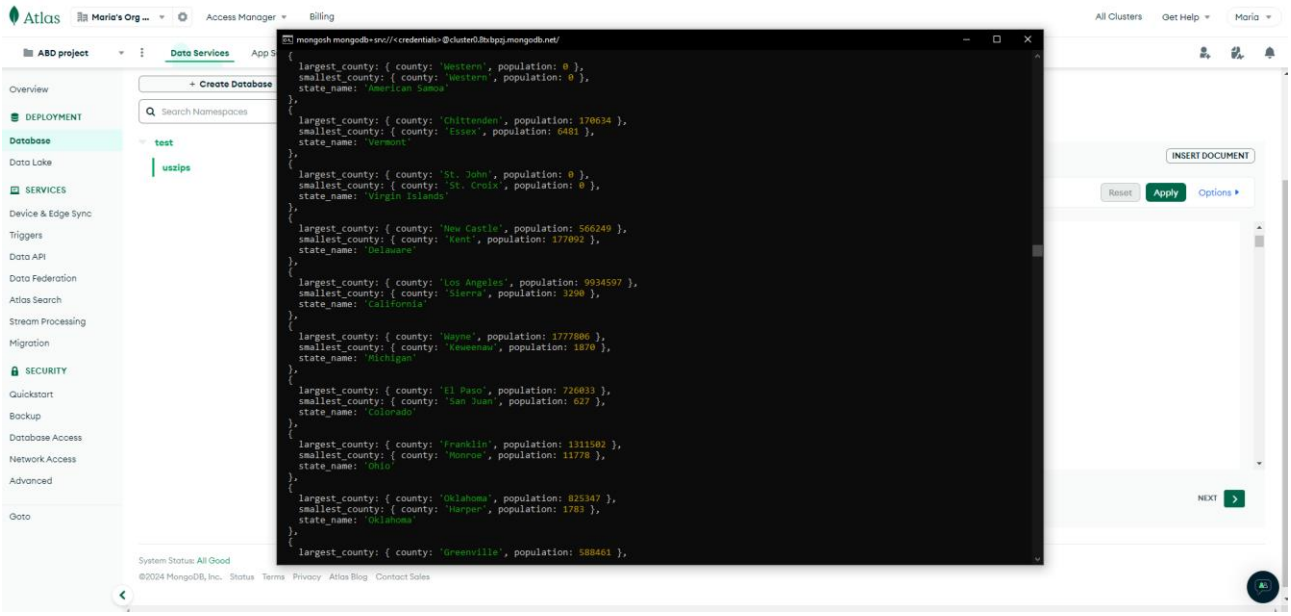

e) Obtineti cele mai apropiate 10 zip-uri de la unul dintre reperele din Chicago, Turnul Willis, situat la coordonatele 41.878876, -87.635918

Pentru a putea folosi capabilitatile geospatiale puse la dispozitie de MongoDB, am adaugat o noua proprietate GeoJSON pentru fiecare document, de forma: "location": { "type": "Point", "coordinates": [-87.6375, 41.8822] }

Am obtinut acest lucru folosind comanda:

db.uszips.updateMany({}, [{ $set: { location: { type: "Point", coordinates: ["$lng", "$lat"] } } }])

Dupa modificari, documentele vor avea urmatoarea forma:

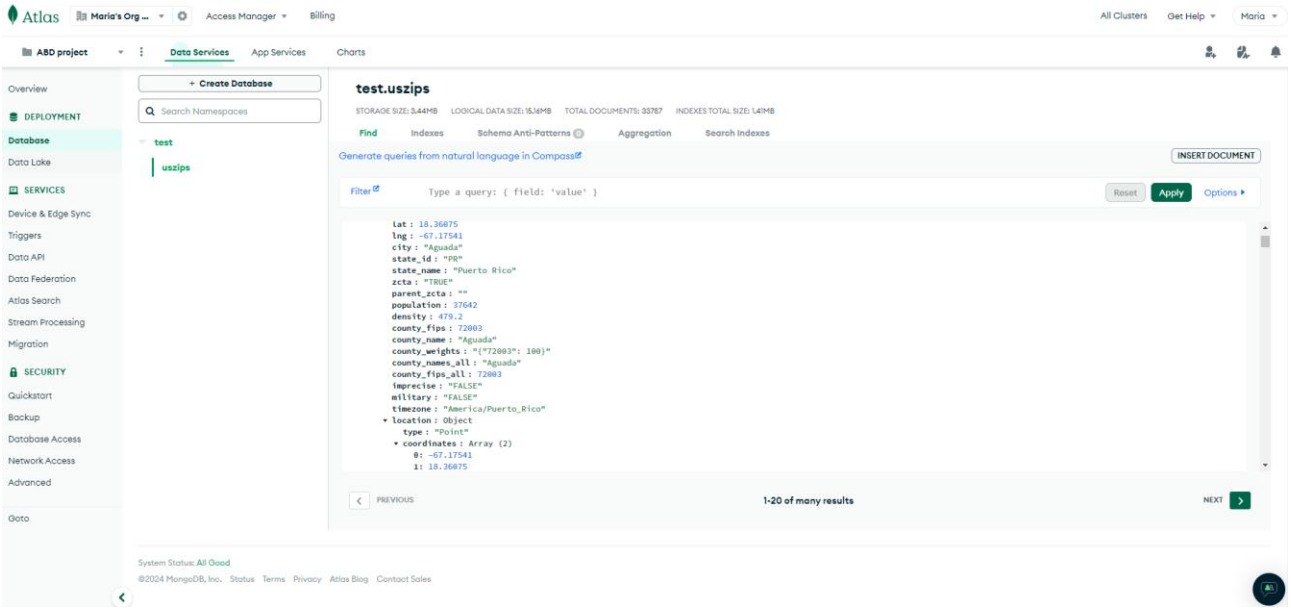

Pentru a cauta cele mai apropiate coduri postale de coordonatele turnului Wiilis, trebuie sa creem un index pentru noua proprietate creata „location”.

db.uszips.createIndex({ location: "2dsphere" })

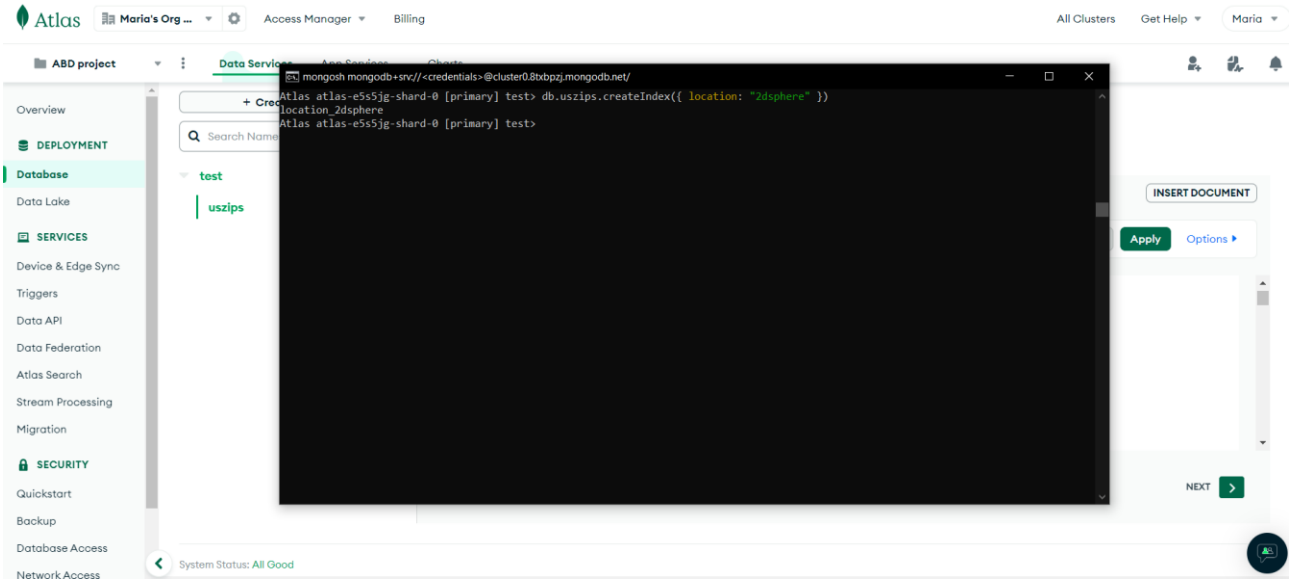

In continuare am utilizat operatorul $near pentru a gasi cele 10 documente necesare, folosind comanda:

db.uszips.find({ location: { $near: { $geometry: { type: "Point", coordinates: [-87.635918, 41.878876] } } } }).limit(10).pretty()

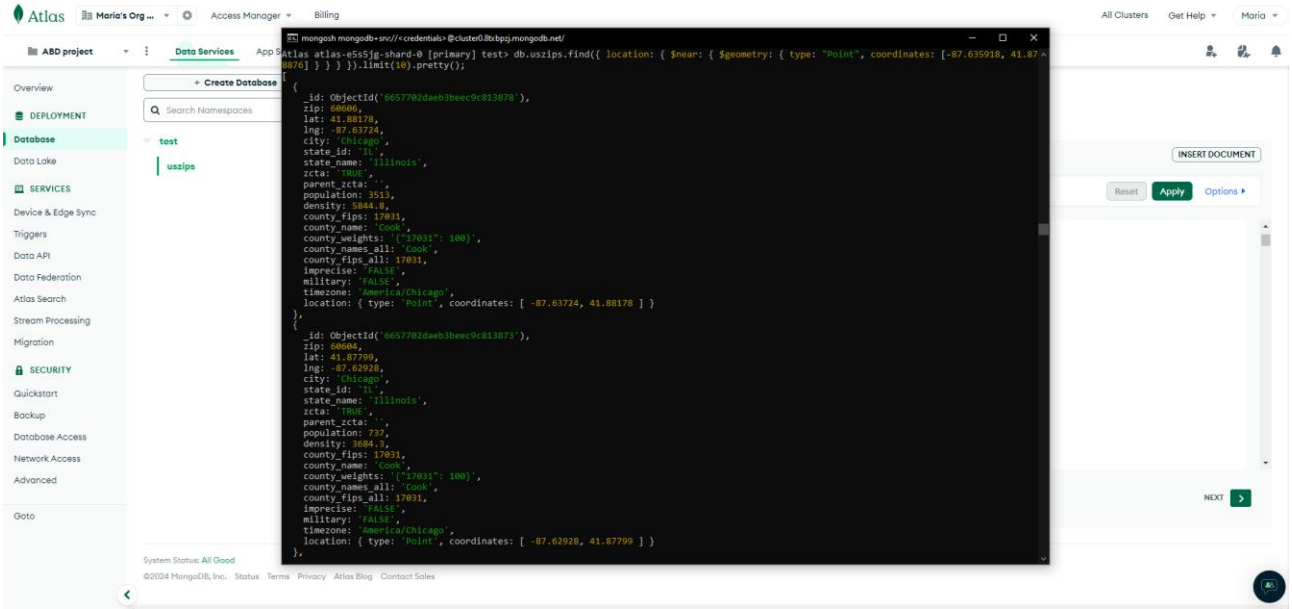

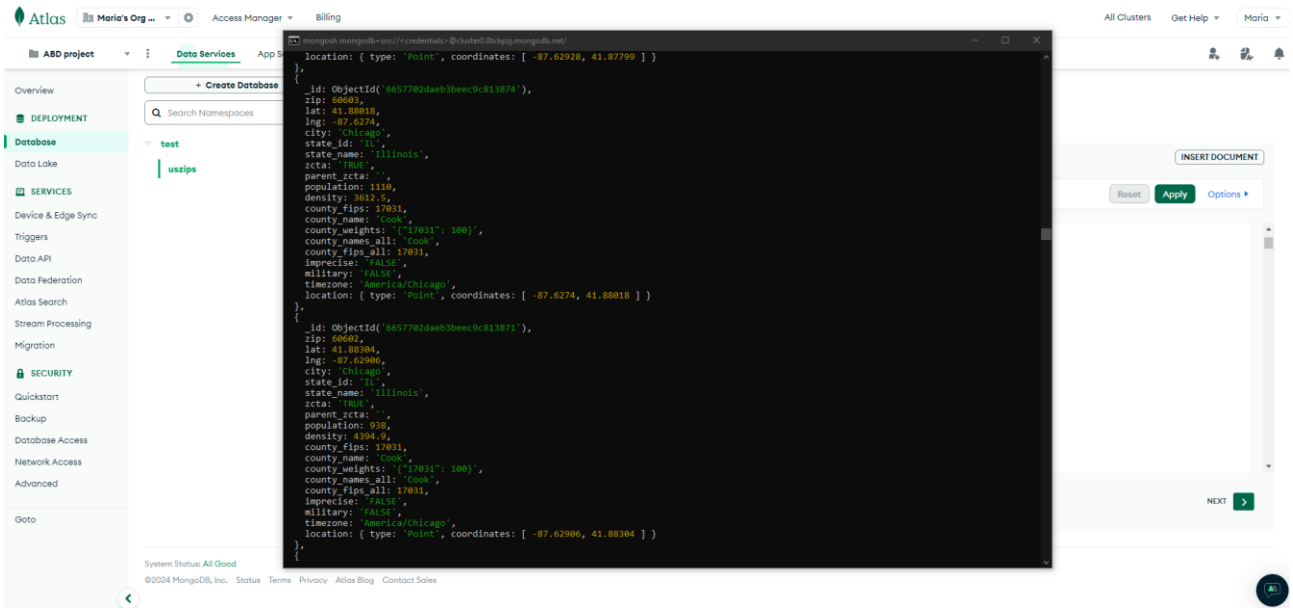
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
Overview
+ Create Database
Search Namespaces
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
Network Access
Advanced
Goto
test
uszips
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good
©2024 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales
mongosh mongodb+srv://<credentials>@cluster0.8xbpcj.mongodb.net/
location: { type: 'Point', coordinates: [-87.62928, 41.87799] }
},
{
_id: ObjectId('6657702daeb3beec9c813874'),
zip: 60603,
lat: 41.88018,
lng: -87.6274,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 1110,
density: 3612.5,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.6274, 41.88018] }
},
{
_id: ObjectId('6657702daeb3beec9c813871'),
zip: 60602,
lat: 41.88304,
lng: -87.62906,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 938,
density: 4394.9,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.62906, 41.88304] }
},
{

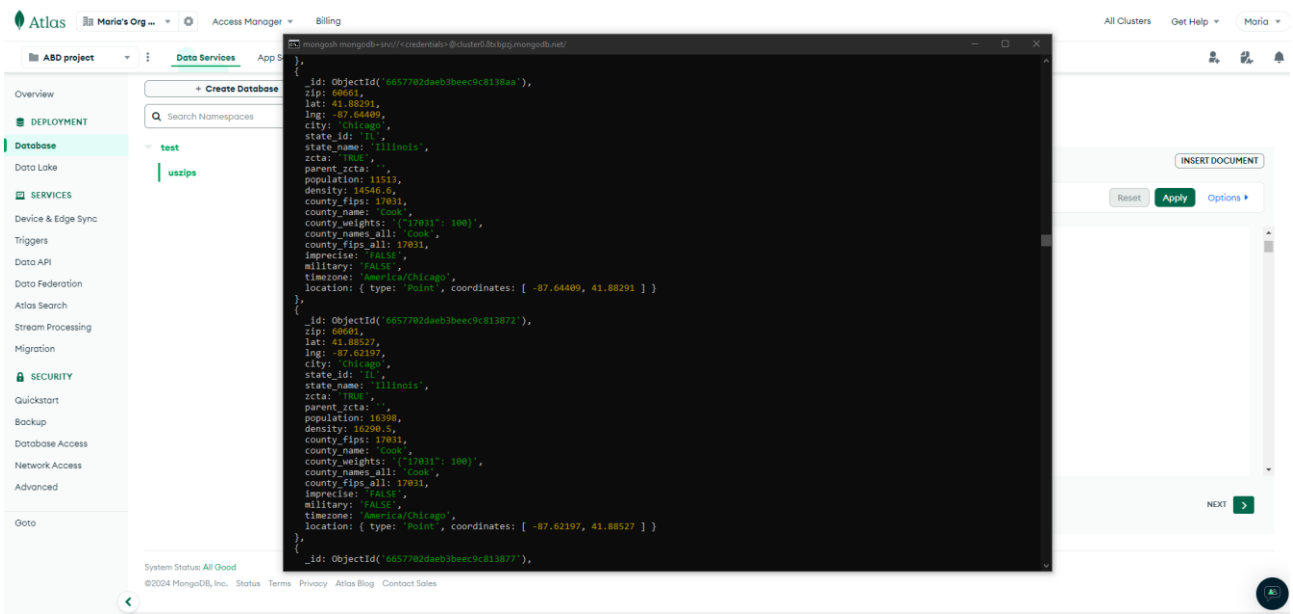
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
Overview
+ Create Database
Search Namespaces
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
Network Access
Advanced
Goto
test
uszips
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good
©2024 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales
mongosh mongodb+srv://<credentials>@cluster0.8xbpcj.mongodb.net/
},
{
_id: ObjectId('6657702daeb3beec9c8138aa'),
zip: 60661,
lat: 41.88291,
lng: -87.64409,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 11513,
density: 14546.6,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.64409, 41.88291] }
},
{
_id: ObjectId('6657702daeb3beec9c813872'),
zip: 60601,
lat: 41.88527,
lng: -87.62197,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 16398,
density: 16290.5,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.62197, 41.88527] }
},
{
_id: ObjectId('6657702daeb3beec9c813877'),

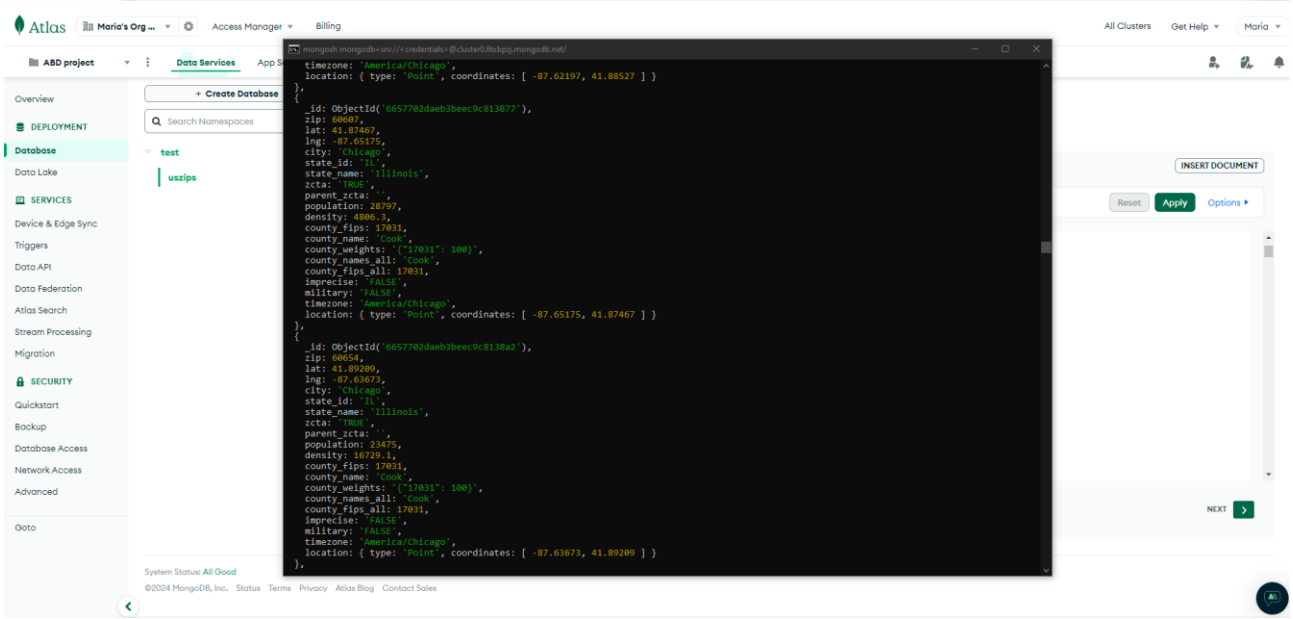
Atlas
Maria's Org ...
Access Manager
Billing
All Clusters
Get Help
Maria
ABD project
Data Services
Overview
+ Create Database
Search Namespaces
DEPLOYMENT
Database
Data Lake
SERVICES
Device & Edge Sync
Triggers
Data API
Data Federation
Atlas Search
Stream Processing
Migration
SECURITY
Quickstart
Backup
Database Access
Network Access
Advanced
Goto
test
uszips
INSERT DOCUMENT
Reset
Apply
Options
NEXT
System Status: All Good
©2024 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales
mongosh mongodb+srv://<credentials>@cluster0.8xbpcj.mongodb.net/
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.62197, 41.88527] }
},
{
_id: ObjectId('6657702daeb3beec9c813877'),
zip: 60607,
lat: 41.87467,
lng: -87.65175,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 28797,
density: 4806.3,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.65175, 41.87467] }
},
{
_id: ObjectId('6657702daeb3beec9c8138a2'),
zip: 60654,
lat: 41.89209,
lng: -87.63673,
city: 'Chicago',
state_id: 'IL',
state_name: 'Illinois',
zcta: 'TRUE',
parent_zcta: '',
population: 23475,
density: 10729.1,
county_fips: 17031,
county_name: 'Cook',
county_weights: '{"17031": 100}',
county_names_all: 'Cook',
county_fips_all: 17031,
imprecise: 'FALSE',
military: 'FALSE',
timezone: 'America/Chicago',
location: { type: 'Point', coordinates: [-87.63673, 41.89209] }
},

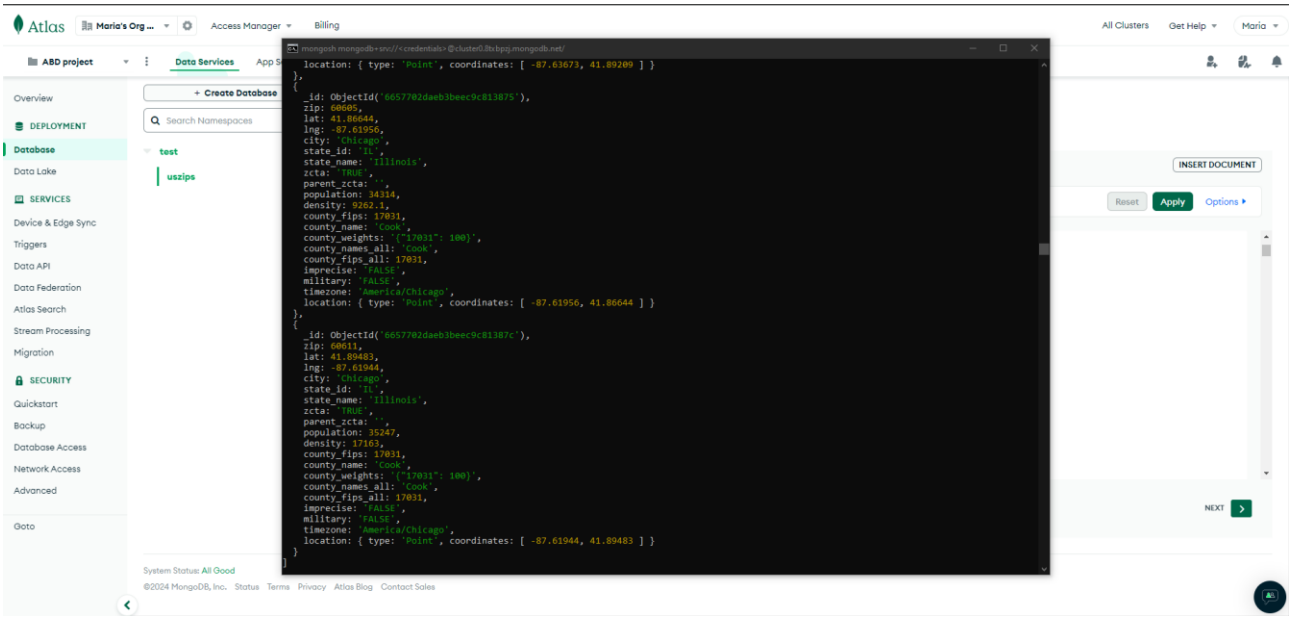

f) Obtineti populatia totala situata intre 50 si 200 km in jurul reperului din New York, Statuia Libertatii, la coordonatele 40.689247, -74.044502

Operatorul $geoWithin obtine documentele care se afla in interiorul unei forme geospatiale, in acest caz forma este definita de o sfera. Aceasta se obtine folosind $centerSphere, care defineste o sfera bazata pe coordonatele centrului si raza. Raza este in radiani, asa ca am impartit distanta (200km) la raza Pamantului (~6378.1 km).

Am folosit comanda:

db.uszips.aggregate([{ $match: { location: { $geoWithin: { $centerSphere: [[-74.044502, 40.689247], 200 / 6378.1] } } } }, { $match: { location: { $not: { $geoWithin: { $centerSphere: [[-74.044502, 40.689247], 50 / 6378.1] } } } } }, { $group: { _id: null, total_population: { $sum: "$population" } } }]).pretty()